

PRIYANKA MOORTHY

408-480-8833 | priyankamoorthycit@gmail.com | [LinkedIn](#) | [GitHub](#)

AI Engineer with 5+ years of professional software and AI engineering experience, specializing in multi-agent orchestration, grounded retrieval, and production LLM evaluation for enterprise healthcare systems.

EXPERIENCE

AI Engineer · Espercare LLC San Jose, CA | Jul 2023 – Present

Serve as the primary technical owner and client-facing AI lead for a payment integrity team at a health insurance client, partnering directly with the client's CTO to identify automation opportunities, scope new AI initiatives, and align technical roadmaps with business priorities.

- Architected and own end-to-end an **agentic reasoning and routing** system (router-orchestrated multi-agent architecture) serving four cross-functional operations teams from a single platform; designed intent-classification routing to five specialized sub-agents (FAQ, exclusions, outreach, letter templates, appeals), achieving **96% task success rate** and **98% out-of-scope rejection rate** at **2–3K queries/day**.
- Own and operate a revenue-code assignment service built on grounded retrieval with **optimized document chunking** over a 1.8M-document corpus; the system handles **~14K line items per day** across multiple claims, replacing ~8 hours of daily manual coding and achieving **98% automation**, with low-confidence cases automatically escalated to human review.
- Designed and own a structured data-extraction pipeline for UB-04, HCFA-1500, and itemized bill documents, previously a fully manual claims-assistant workflow; auto-extracts **96% of records** at high confidence, with residual cases routed to review; drove a **27% accuracy lift** through targeted domain-specific fine-tuning.
- Built and maintain a continuous evaluation harness on **Vertex AI** covering pairwise preference scoring, human feedback loops, and regression gating across all agents and pipelines; established an **LLM-as-judge framework** for automated output QA, enabling confident, gated rollouts and preventing quality regressions at scale.
- Collaborate cross-functionally with claims, operations, and engineering teams to gather workflow requirements, translate business logic into agent behaviors, and ensure consistent adoption and feedback loops across teams.

Software Engineer · RedBlackTree Tech India | Jan 2019 – May 2020

- Owned and shipped an OCR + CNN classification pipeline to extract provisions from lease documents, reducing processing time by **82%** at **0.92 AUC**.
- Designed and owned an async processing system on **RabbitMQ and Celery** to distribute model inference across workers and control concurrency under production load.

PERSONAL PROJECTS

Output Observatory · Interactive LLM Evaluation Learning Platform

- Designed and built an interactive learning platform for LLM evaluation covering RAG evaluation, agent evaluation, finetuned model testing, and production monitoring across 6 tracks and 24 hands-on lessons targeting ML engineers and practitioners.
- Integrated industry-standard evaluation frameworks (RAGAS, DeepEval, LangSmith, Guardrails AI, Promptfoo) to surface real metrics including faithfulness, context precision, hallucination detection, and trajectory efficiency, giving learners hands-on exposure to production eval workflows.

Prior Authorization Orchestration · Multi-Agent PA Automation with Claude & MCP

- Built a multi-agent Prior Authorization orchestrator using Claude and custom MCP servers to automate payer criteria retrieval and clinical decision-making, with confidence-gated human-in-the-loop escalation.
- Built a full-stack integration layer connecting FHIR patient data and payer PDF policy systems via enterprise APIs (MCP tool interfaces), demonstrating how agentic AI workflows integrate with external enterprise IT systems in regulated healthcare settings.

TECHNICAL SKILLS

AI Systems: Multi-agent orchestration, RAG / grounded retrieval, function & tool calling, LLM evaluation (pairwise preference, LLM-as-judge, human feedback), prompt & system design, model fine-tuning

Developer Tools: Claude Code, MCP (Model Context Protocol), Claude Skills, GitHub Copilot

ML & Modeling: Deep learning, NLP, reinforcement learning, computer vision

Infrastructure & MLOps: GCP (Vertex AI, Cloud Run), AWS, Docker, Kubernetes, RabbitMQ, Celery, async job systems, model monitoring

Frameworks & Libraries: PyTorch, TensorFlow, Transformers, LangChain, LlamaIndex, Scikit-learn, OpenCV

Languages: Python, SQL, Java, C++, JavaScript, Shell

CERTIFICATIONS

Claude Code in Action · Anthropic Education · May 2026 · verify

EDUCATION

San Jose State University · M.S., Artificial Intelligence · San Jose, CA · Aug 2021 – May 2023

Anna University · Integrated M.Sc., Software Systems · India · Aug 2014 – May 2019